

Shubham Chougale

9156719927 | shubham.chougale001@gmail.com | portfolio | LinkedIn | GitHub | Pune, Maharashtra

SUMMARY

Production AI/ML Engineer with 2+ years architecting and deploying LLM-powered systems serving 10,000+ daily users. Specialized in generative AI, multi-agent orchestration, RAG pipelines, Voice AI, and Claude integration. Delivered measurable impact across every role: 92% CV accuracy, 40% query resolution improvement, 65% reduction in manual intervention, and 99.2% production uptime. Current focus on MCP server design, agentic system engineering, and Claude-native application development.

TECHNICAL SKILLS

Core AI/ML: Large Language Models (LLMs), Generative AI, Claude API (Sonnet, Opus), RAG Pipelines, LangChain, LangGraph, CrewAI, Multi-Agent Systems, Agentic AI, Prompt Engineering, Tool Calling, Function Calling
Model Frameworks & Tools: Claude, OpenAI, Gemini, MCP, FastMCP, Langfuse, LLM Fine-Tuning, Windsurf IDE, Groq Cloud
Voice AI & Real-Time Systems: Pipecat, Deepgram, ElevenLabs, Twilio, Plivo, Speech Synthesis, Real-Time Voice Pipelines
Vector Search & Retrieval: ChromaDB, FAISS, Pinecone, Embedding Models, Semantic Search, Vector Databases
ML & Computer Vision: YOLOv8, CoreML, OpenCV, CNNs, ResNet, Feature Engineering, Model Deployment
Backend & Cloud: Python, FastAPI, Django, REST APIs, Microservices, MySQL, PostgreSQL, Snowflake, AWS (Lambda, Step Functions, Glue, S3, EC2, IAM, CloudWatch, SNS), Docker, Terraform, Jenkins, Git
Knowledge Graphs: Neo4j, Ontology Design, Graph-Based Data Modeling

EXPERIENCE

Software Engineer — Gen AI & Python | Coditas Jan 2026 – Present
– Building an AI-powered workflow automation platform processing 100+ automations daily using Claude Sonnet and Claude Opus, delivering a 92% data extraction accuracy via RAG-augmented prompt chains (up from 78%).
– Architected Model Context Protocol (MCP) server integration with 10+ custom domain-specific MCP servers, expanding system capabilities 5x and enabling any MCP-compatible LLM to interface with the platform.
– Reduced workflow execution latency by 20% and token costs by 25% through intelligent LLM response caching and dynamic model selection across Claude, OpenAI, and Groq backends.
– Shipped 5+ major features (workflow composition, debugging tools, model switching, result export) with 95%+ test coverage and zero critical regressions; resolved 15+ LLM integration bugs improving response consistency by 18%.
AI/ML Engineer | 64-Squares LLC Mar 2024 – Oct 2025
– Deployed YOLOv8-based defect detection model on iOS via CoreML achieving 92% accuracy at 30 FPS, processing 100+ property inspections daily with <200ms inference latency across 500+ properties.
– Constructed a production RAG system (OpenAI + FAISS/ChromaDB) reducing tenant query resolution time by 40% while cutting API costs by 30%, serving 2,000+ monthly interactions with 95%+ response accuracy.
– Architected multi-agent workflows using LangGraph and CrewAI managing 50+ concurrent agent instances, reducing manual real estate task intervention by 65% across automated pipelines.
– Optimized AI inference servers (FastAPI) for 99.2% uptime and sub-500ms latency under 10,000+ concurrent users; built Snowflake/SQL data pipelines cutting model training time by 30%.

EDUCATION

MIT-ADT University	Master of Computer Applications (Data Science) — CGPA : 8.58	2021 – 2023
Shivaji University	B.Sc. in Computer Science — CGPA : 8.86	2018 – 2021

PROJECTS

Sieve | Claude, Neo4j, Ontology, Knowledge Graphs, CRM Integration, Slack Webhooks, BoldSign, Google Drive 2026
– Designed a Claim Dependency Graph and integrated Neo4j for graph-based storage/visualization of claim-to-vocabulary relationships, building an interactive claim/concept graph view and fixing a data persistence bug that silently dropped large responses.
– Automated NDA workflows across BoldSign, Google Drive, and CRM, including centralized status mapping, automatic company-folder creation, and re-filing of previously unmatched NDAs, with Slack webhook alerts on sync failures.
SECRISA — AI Manufacturing Chatbot | LangChain, RAG, ChromaDB, Tool Calling, FastAPI, Python 2025
– Launched a domain-specific manufacturing chatbot (RAG + LangChain) for pricing, SKU lookup, budgeting, and product categorization, reducing manual lookup time by 70% across 200+ daily operations.
– Implemented Tool Calling and Function Calling APIs with dynamic agent routing, enabling accurate multi-turn query execution across heterogeneous product and financial datasets.
ViziSmart — AI Property & Voice Assistant | LangGraph, LangChain, RAG, OpenAI, ChromaDB, Django, AWS 2024
– Crafted a voice-enabled real estate assistant with LangGraph multi-agent routing for property info retrieval, ticket management, and tenant operations, achieving 88%+ multi-turn response relevance.
– Deployed on AWS (EC2, S3, Lambda) with CloudWatch monitoring; accelerated LLM-based content generation by 40% through async orchestration and prompt optimization, maintaining 99.5% uptime.

CERTIFICATIONS

Claude 101 | Claude Code 101 | Claude Code in Action | Introduction to Agent Skills | Introduction to Model Context Protocol — Anthropic
Multi AI Agent Systems with CrewAI — DeepLearning.AI
Large Language Models with Semantic Search — Cohere / DeepLearning.AI
SnowPro Core Certification — Snowflake